

SYNTHETIC ESTIMATE — 200 heavy ML workloads (not a live Kubernetes run)

