



Generated feature distributions — 200 synthetic ML workloads

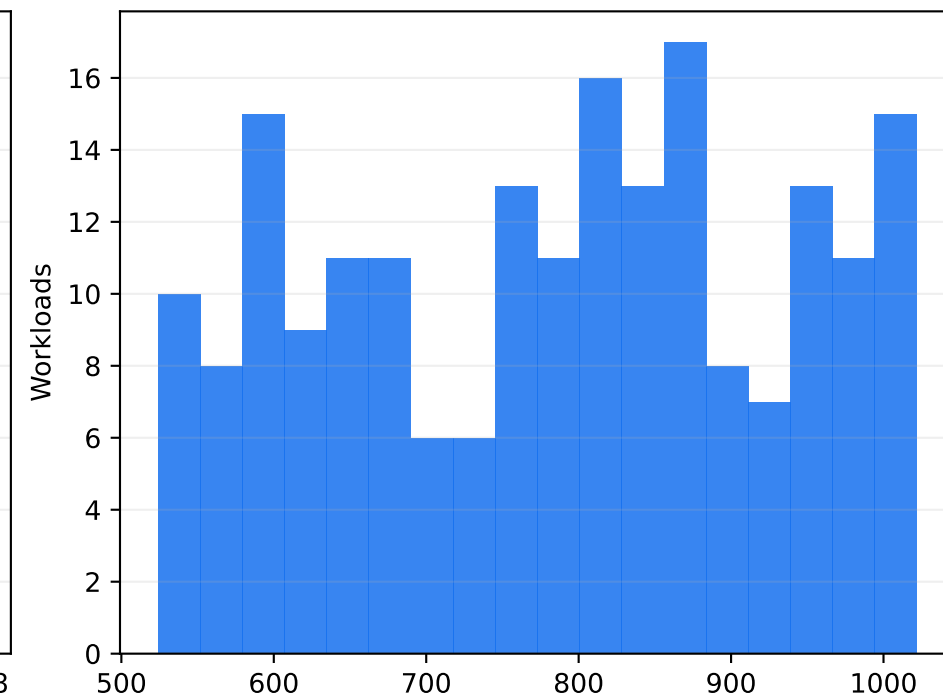
T: estimated duration



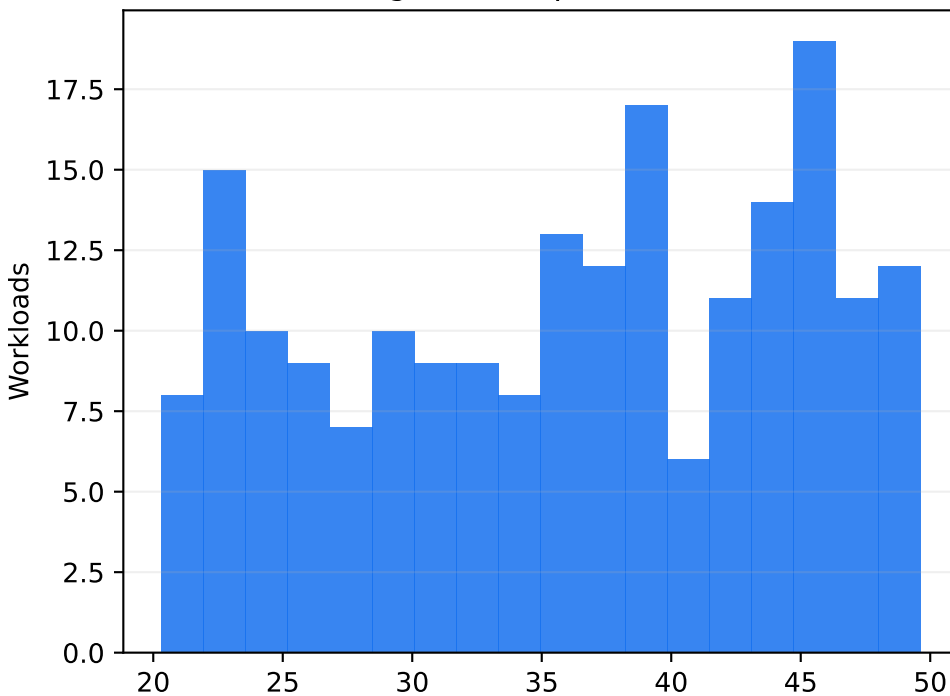
R: loss-reduction rate



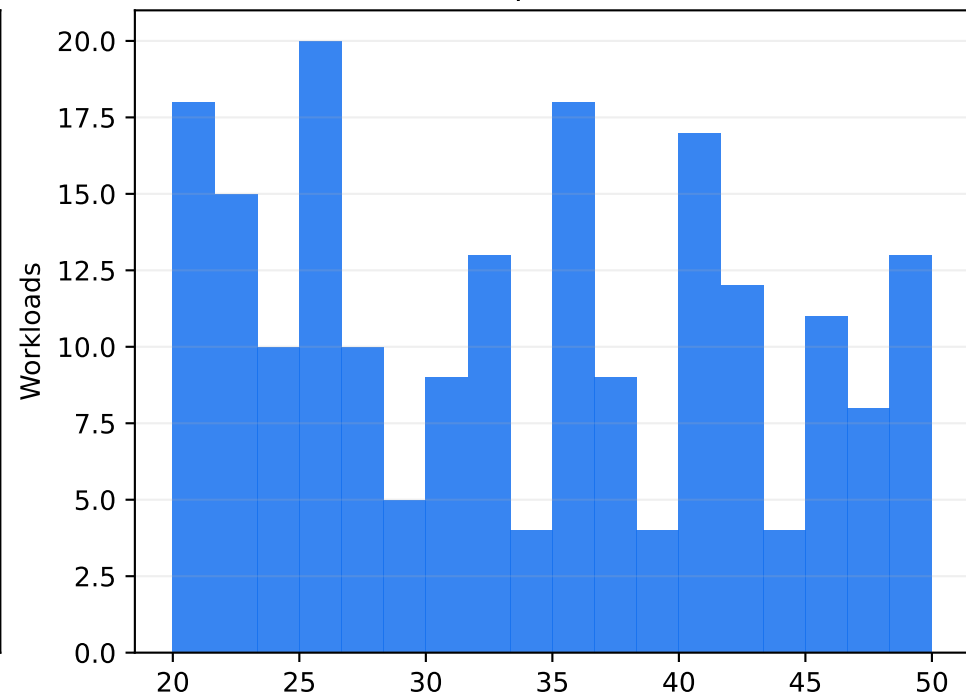
M: matrix dimension



G: gradient update (MiB)



C: checkpoint interval



P: model partitions

